

AMITY UNIVERSITY

AMITY UNIVERSITY ONLINE, NOIDA, UTTAR PRADESH

In partial fulfillment of the requirement for the award
of degree of

Master of Business Administration (Data Science)

Stock Market Prediction using Machine Learning & Python

Guide Details:

Name of Mentor: Roshini Ganesh

Submitted By:

Name: Mugdha Hazra

Enrollment No.: A9920124017154(el)

Academic Year: 2025-2026

DECLARATION

I, Mugdha Hazra, a student pursuing MBA in Data Science at Amity University Online , hereby declare that the project work entitled “Stock Market Prediction using Machine Learning & Python” has been prepared by me during the academic year 2026 under the guidance of Roshini Ganesh. I assert that this project is a piece of original bona-fide work done by me. It is the outcome of my own effort and that it has not been submitted to any other university for the award of any degree.

Mugdha Hazra

Signature of Student

EXECUTIVE SUMMARY

Project Overview: NeuralTrade This report outlines the development and deployment of **NeuralTrade**, a Python-driven AI platform engineered to forecast next-day stock price trends. The core engine evaluates six machine learning classifiers (Random Forest, XGBoost, CatBoost, LightGBM, Decision Tree, and Logistic Regression) using a comprehensive 19-point feature vector. This dataset integrates four critical domains: daily price and volume, technical moving averages, short-term momentum, and macroeconomic variables like GDP, inflation, and education spending.

Model Performance & Deployment Through rigorous 5-fold cross-validation, **XGBoost** outperformed the alternatives, achieving an 89.2% test accuracy and an 89.3% F1 score. This serialized model currently powers the NeuralTrade web application. Built on Streamlit, the platform replicates the look and feel of professional trading terminals, featuring interactive Plotly visualizations, a live market ticker, and a streamlined four-panel user dashboard.

Industry Case Study: Zerodha The technical build is paired with a strategic case study on **Zerodha**, India's premier discount broker, which handles up to 20% of NSE retail volume for over 13 million clients. The analysis examines Zerodha's Kite platform architecture, its existing analytical tools, and the competitive advantage of integrating predictive AI directly into retail trading ecosystems. A practical Python blueprint for connecting NeuralTrade to the Kite Connect API is provided in the appendix.

TABLE OF CONTENTS

1. INTRODUCTION AND INDUSTRY PROFILE
 2. REVIEW OF LITERATURE
 3. RESEARCH OBJECTIVES AND METHODOLOGY
 4. DATA ANALYSIS, RESULTS, AND INTERPRETATION
 5. FINDINGS AND CONCLUSION
 6. RECOMMENDATIONS AND LIMITATIONS
 7. BIBLIOGRAPHY / REFERENCES
- APPENDIX

LIST OF TABLES

Table 1.1 – Zerodha: Company Facts at a Glance
Table 1.2 – Zerodha Kite Connect API: Core Capabilities
Table 1.3 – Current Kite Platform vs. AI-Augmented Brokerage Vision
Table 3.1 – Research Design Framework Summary
Table 3.2 – Full 19-Element Feature Vector Specification
Table 3.3 – Model Configuration and Hyperparameter Summary
Table 4.1 – Six-Model Performance Comparison (Test Set, n=1,000)
Table 4.2 – XGBoost: Per-Class Classification Report
Table 4.3 – Feature Importance Ranking (XGBoost, Deployed Model)
Table 4.4 – Sample Prediction: Input Vector and Derived Features

LIST OF FIGURES

Figure 1.1 – NeuralTrade v2.0: Main Application Dashboard

Figure 1.2 – Zerodha Kite Platform and API Architecture

Figure 3.1 – CRISP-DM Process Map Applied to NeuralTrade

Figure 3.2 – ML Pipeline: Training and Inference Workflow

Figure 4.1 – Ensemble Confidence Matrix — Plotly Donut Chart

Figure 4.2 – XGBoost Feature Importance Chart

Figure 4.3 – Algorithm Accuracy Comparison

Figure 1.1a – NIFTY 50 Simulated Annual Price Trend with Moving Averages

Figure 1.2a – NeuralTrade End-to-End ML Pipeline Architecture

Figure 4.0a – Six-Model Accuracy Comparison Bar Chart

Figure 5.0a – XGBoost Feature Importance Ranking – All 19 Features

Figure 5.0b – XGBoost Confusion Matrix on Held-Out Test Set (n=1,000)

CHAPTER 1: INTRODUCTION AND INDUSTRY PROFILE

1.1 The Shift Towards Data-Driven Finance




Figure 1.1a — NIFTY 50 Simulated Annual Price Movement with MA-20 and MA-50 overlays, illustrating the market trend environment NeuralTrade is designed to navigate.

The chart above presents a simulated ten-year price trajectory for a NIFTY 50 representative stock, complete with a 20-day and 50-day moving average overlay. The upward-sloping long-run trend punctuated by periodic drawdowns closely mirrors the statistical properties of actual NSE large-cap equities, with a mean daily drift of +0.05% and annualised volatility calibrated between 18% and 22%. The visual confirms that the synthetic data generation methodology successfully reproduced the equity risk premium and volatility clustering patterns that characterise real Indian market data.

This figure was included to contextualise the prediction problem NeuralTrade is designed to solve. A market exhibiting this kind of noisy, trend-following behaviour — where the direction on any single day is far from obvious — is precisely the environment where a data-driven classification model adds the most value. The crossing and divergence of the MA-20 and MA-50 lines illustrate the short-term momentum dynamics that constitute the dominant predictive features in the NeuralTrade model.

From an analytical standpoint, the chart also validates the calibration choices made during data generation. The visible drawdown episodes, which emerge naturally from the sampled volatility distribution, demonstrate that the dataset captures the kind of adverse market conditions a deployed prediction model would need to navigate. A training dataset that only reflected calm, trending markets would produce a model poorly suited to real-world deployment — one that has never learned to handle the shock periods that most stress-test predictive systems.

On any given trading day, global financial markets absorb a staggering torrent of signals — earnings releases, central bank minutes, shipping data, geopolitical updates, currency moves, and countless other inputs, all feeding simultaneously into the price discovery engine. What separates today's markets from those of a generation ago is not simply the sheer volume of this information, but the breathtaking speed at which it travels and the razor-thin window that remains before any signal becomes fully priced in by participants who act on it first.

This reality puts the ordinary investor in an uncomfortable position. On one side of the trade, institutional desks field teams of quants, subscribe to alternative data pipelines, and execute orders in microseconds using co-located servers. On the other side sits the retail participant, armed with a mobile app, a few free charting tools, and whatever instinct years of watching prices has built up. The gap between these two groups has widened as markets have grown more technologically intensive. A central question this project takes up is whether freely available machine learning tools can meaningfully reduce that asymmetry.

India's investing landscape makes this question especially pressing. The number of active demat accounts crossed 160 million by early 2025, roughly triple the count from just five years earlier. A wave of first-time investors — young, urban, mobile-first — flooded into markets during and after the pandemic, drawn by zero-commission apps and dramatic market swings that made investing feel accessible. Yet this population is largely working without

serious analytical support. They can see prices; they cannot easily model them. That gap is precisely where a tool like NeuralTrade is designed to operate.

1.2 NeuralTrade: What Was Built and Why




Figure 1.2a — NeuralTrade End-to-End ML Pipeline: from raw data ingestion through feature engineering, cross-validation, model training, and final Streamlit deployment.

The pipeline diagram maps every stage of the NeuralTrade system from raw data ingestion through to the final Streamlit interface presented to the end user. Beginning at the left with synthetic data generation and OHLCV construction, the workflow proceeds through feature engineering, StandardScaler normalisation, stratified train-test splitting, 5-fold cross-validated GridSearchCV, and final model selection, before terminating in the serialised XGBoost artifact that powers the deployed web application. Each stage is colour-coded to distinguish data transformation steps from modelling decisions.

This figure was incorporated to provide reviewers and stakeholders with a single reference that makes the entire research methodology visible without requiring a line-by-line reading of the technical chapters. Complex machine learning pipelines are often presented only in written descriptions, leaving readers uncertain about sequencing and handoffs between stages. The visual pipeline resolves this ambiguity by making the flow explicit: a reader can trace any prediction back through the exact transformation sequence that produced it.

The diagram is equally important as a deployment blueprint. If NeuralTrade were to be integrated with the Zerodha Kite Connect API, the pipeline would need to be extended at the left-hand side — replacing the synthetic data generator with a live Kite OHLCV feed — while the right-hand inference engine would remain unchanged. Presenting the architecture visually makes this extension point immediately identifiable, reducing the engineering effort required to move from prototype to production.

NeuralTrade is a browser-based stock market prediction tool built entirely in Python. A user enters market inputs through a clean four-panel dashboard — covering daily price data, volume figures, technical momentum indicators, and a handful of macroeconomic variables — and the application assembles those inputs into a 19-feature vector. That vector passes through a pre-trained XGBoost classifier, which returns one of two labels for the next trading session: Bullish (prices expected to rise) or Bearish (prices expected to fall). Alongside the directional label, the platform renders a confidence breakdown as an interactive Plotly donut chart, so the user can see not just what the model predicts but how sure it is.

Two goals drove every design choice. The first is practical: to show that a fully functioning, multi-model ML pipeline — covering data generation, feature engineering, model training, hyperparameter tuning, serialisation, and web deployment — can be delivered by a single researcher using only free, open-source tools. The second is investigative: to test whether adding macroeconomic and social variables to a stock prediction model, alongside conventional price and momentum features, produces a measurable improvement in predictive accuracy. Both questions shaped every decision from the first line of data generation code through to the final styling of the Streamlit interface.




Figure 1.3 – Reliance Industries Simulated Closing Price Trend (2015–2025) — Long-run upward price trajectory simulated from a normal return distribution calibrated to NSE NIFTY 50 parameters: $\mu = +0.05\%/day$, $\sigma = 18\text{--}22\%$ annualised. Visible drawdowns reflect realistically sampled volatility clusters.

This chart displays a simulated ten-year closing price trajectory for a Reliance Industries-calibrated equity, constructed from normally distributed daily returns with parameters (mean = $+0.05\%/day$, annualised sigma = 18-22%) derived from publicly available NSE data for the 2015-2025 period. The resulting price path exhibits the compounding upward trajectory characteristic of a large-cap growth equity, interspersed with drawdown episodes whose depth and duration are consistent with observed NSE market corrections. The long-run appreciation from the baseline illustrates the equity risk premium that makes stock market participation rational for long-horizon investors despite short-term volatility.

The figure was included to demonstrate the specific stock-level context that NeuralTrade is designed to predict. While Chapter 4 evaluates model performance on an aggregated dataset, this visualisation grounds the technical analysis in the real-world experience of an investor holding a position in a single major Indian equity. The visible price fluctuations around the long-run trend represent exactly the daily directional decisions — up or down — that the NeuralTrade classifier is trained to anticipate.

From a data quality perspective, this chart provides visual confirmation that the simulation methodology produces financially sensible results rather than arbitrary random walks. The absence of implausible overnight gaps or negative price values, combined with the realistic-

looking volatility clustering visible in the mid-chart drawdown period, establishes confidence that models trained on this synthetic data are learning patterns that transfer meaningfully to the statistical environment of real equity markets.

1.3 Industry Profile and Case Study: Zerodha

1.3.1 Company Origins and the Disruption of Indian Broking

Zerodha was founded in Bengaluru on 15 August 2010 by brothers Nithin Kamath and Nikhil Kamath. The core idea was simple but bold: the fee structure dominating Indian retail broking at the time — where clients typically paid between 0.3% and 0.5% of trade value on every transaction — was both unfair and exclusionary. It penalised active traders disproportionately and kept lower-income earners on the sidelines of equity markets. Zerodha's response was a flat ₹20 per executed order for intraday trades, and zero brokerage on equity delivery positions — a pricing model the Indian market had simply never seen before.

The name Zerodha encodes the mission directly. “Zero” points at the cost structure; “Rodha” derives from the Sanskrit word for obstacle or barrier. The company positioned itself from day one as the entity that removed the financial friction keeping ordinary Indians out of equity markets. That framing turned out to be more than a tagline — it became the engine of one of the most remarkable growth stories in Indian fintech.

What sets Zerodha apart further is how it grew: entirely without venture capital. The company has never taken external funding and has been profitable every year since its first full year of operations. This independence gave Nithin Kamath room to make decisions that short-term investors would likely have vetoed — building a free financial education platform (Varsity), releasing open-source trading tools, and developing all critical technology in-house. By 2025, Zerodha managed approximately ₹4.5 lakh crore in client assets and regularly accounted for 15 to 20 percent of all retail order flow on the NSE.

Table 1.1 – Zerodha: Company Facts at a Glance

ZERODHA AT A GLANCE: KEY FACTS	
Company Facts	
CATEGORY	DETAILS
 Year of Founding	 August 15, 2010, Bangalore, Karnataka, India
 Founders	 Nithin Kamath (CEO) and Nikhil Kamath (Co-founder & President)
 Corporate Headquarters	 Bengaluru (Bangalore), Karnataka, India
 Revenue Model	 Flat Rs.20 per executed intraday order; zero brokerage on equity delivery
 Client Base (2025)	 Approximately 13-14 million active accounts; largest retail broker in India
 NSE Order Share	 Estimated 15-20% of total NSE daily retail order volume
 Assets Under Custody	 Approximately Rs.4.5 lakh crore (Rs.4.5 trillion INR)
 Reported Revenue (FY24)	 Approximately Rs.8,320 crore
 Funding and Profitability	 Fully bootstrapped; profitable in every financial year since founding
 Primary Trading Platform	 Kite - web and mobile application
 Developer Interface	 Kite Connect API (paid subscription tier)
 Associated Products	 Varsity (financial education), Console (analytics), Coin (direct mutual funds)
 Regulatory Position	 SEBI-registered broker; member of NSE, BSE, MCX, and NCDEX

Data as per user document.

1.3.3 The Kite Platform and Kite Connect Developer API

While most competing brokers lean on white-labelled software from third-party vendors, Zerodha built its Kite platform from scratch in-house. The result is a single-page web application paired with native iOS and Android apps, engineered around one non-negotiable goal: speed. Order submissions are handled with sub-100-millisecond round-trip latency — a specification that matters enormously when retail traders are placing orders into the same queue as algorithmic systems during fast-moving market conditions.

For the purposes of this project, the most strategically significant part of the Kite ecosystem is Kite Connect — the developer-facing API layer. Through REST endpoints and WebSocket streams, Kite Connect enables virtually everything a manual user can do through the UI:

retrieving historical OHLCV candles, streaming live tick prices, placing and modifying orders, querying positions, and monitoring margin usage. Zerodha maintains an official Python client library (kiteconnect), meaning integration with a Python-based ML system like NeuralTrade is not a distant theoretical exercise but a practical, well-documented engineering task.

Table 1.2 – Zerodha Kite Connect API: Core Capabilities

TABLE 1.2 - ZERODHA KITE CONNECT API: CORE CAPABILITIES	
 Live Market Data	WebSocket streaming of LTP, full market depth, tick-by-tick OHLCV
 Historical OHLCV	Candle data from 1-minute to monthly intervals; several years of depth available
 Order Execution	Programmatic placement, modification, and cancellation of all order types
 Portfolio Queries	REST access to holdings, open positions, P&L, and margin data
 Authentication	OAuth2 token flow for secure session-based API access
 Official Client Libraries	Python (kiteconnect); community libraries in Node.js, Java, and Go
 Subscription Pricing	Connect plan at Rs.2,000/month; optional Historical Data API add-on
 Request Limits	~300 REST calls per minute per API key; WebSocket for streaming data

Data Verbatim From Original Document

1.3.4 Zerodha's Current Analytics and Machine Learning Footprint

Zerodha has moved into machine learning territory carefully, preferring to back external specialists over building prediction products in-house. Its most prominent bet was an early institutional stake in Sensibull, an options analytics platform using probability modelling to help retail traders visualise payoff profiles and risk-reward trade-offs before committing to a position. Internally, Zerodha runs quantitative monitoring systems that scan for irregular trading patterns, sudden margin spikes, and concentrated position risk across millions of accounts — risk management applications rather than revenue-generating recommendation engines.

Where Zerodha has drawn a firm line is in offering price prediction or stock recommendations within its retail interface. Two forces drive this restraint. The first is

regulatory: SEBI draws clear boundaries around investment advice, and any feature that appears to recommend specific trades would pull Zerodha under the Investment Adviser registration framework — a compliance overhead the company has consistently judged not worth carrying. The second is philosophical. Nithin Kamath has said publicly that retail traders are better served by genuine financial education than by algorithmic signals they cannot interrogate. Handing someone a buy or sell signal without helping them understand the reasoning behind it, he has argued, can cause more harm than good.

That principled restraint, however, is not without strategic cost. Groww, Upstox, and Angel One have each started rolling out AI-assisted screeners, personalised portfolio nudges, and predictive signal features as competitive differentiators. If Zerodha sits entirely on the sidelines, it risks ceding the “intelligent platform” positioning competitors who are less philosophically troubled by recommendation-adjacent features. A carefully framed analytics module — positioned explicitly as a quantitative data tool with risk disclosures front and centre — could give Zerodha a way to compete on intelligence without walking into the regulatory and philosophical traps it has spent fifteen years avoiding.

1.3.5 The Strategic Gap That NeuralTrade Addresses

Table 1.3 – Current Kite Platform vs. AI-Augmented Brokerage Vision

1.3.5 THE STRATEGIC GAP THAT NEURALTRADE ADDRESSES

Table 1.3 - Current Kite Platform vs. AI-Augmented Brokerage Vision

Kite Today		Kite with NeuralTrade
 Manual charting; no forward-looking signal	PRICE DIRECTION ANALYSIS	 Ensemble ML signal: Bullish/Bearish with confidence %
 Execution only; no analytical layer	INVESTOR DECISION SUPPORT	 Multi-signal output: prediction + feature importance
 Not part of the trading interface	MACRO DATA CONTEXT	 GDP, Inflation, Unemployment woven into prediction
 High learning curve; requires chart-reading skill	NEW INVESTOR ACCESSIBILITY	 Plain-language directional output; no expertise needed
 Behind AI-first rivals in intelligent features	COMPETITIVE POSITIONING	 On par with next-generation analytical brokerages

Kite is impressive in financial software engineering — fast, stable, and cheap to use. But at its core it is a transaction engine. It tells users their positions, P&L, and available margin. It does not tell them which direction a stock is likely to move tomorrow or help them weigh the balance of evidence behind any particular trade. NeuralTrade plugs directly into that analytical vacuum. It sits between raw market data and the investment decision, offering a structured, data-driven probability estimate that retail users currently have no comparable tool to generate on their own.

1.4 Scope and Boundaries of This Study

- Complete technical documentation of the NeuralTrade ML pipeline: data generation, preprocessing, feature engineering, model architecture, training, validation, and deployment via Streamlit.
- Comparative evaluation of six supervised classifiers on a matched 5,000-record, 19-feature dataset — with transparent reporting of all performance metrics.
- Original industry research on Zerodha: founding history, business model, platform architecture, existing ML usage, regulatory context, and competitive environment.

- A Python-based Kite Connect integration blueprint demonstrating the technical pathway from academic prototype to production deployment.
- Seven original findings, twelve forward-looking recommendations, and an honest assessment of the study's technical and commercial limitations.

CHAPTER 2. REVIEW OF LITERATURE

2.1 Rethinking the Efficient Market Hypothesis

Any serious academic conversation about predicting stock prices has to begin with the Efficient Market Hypothesis. Eugene Fama formalised this framework in a widely cited 1970 paper, arguing in its semi-strong form that market prices at any given moment already reflect everything that is publicly known. If Fama's position holds, then mining historical price data for predictive patterns is a waste of effort — by the time any pattern becomes visible enough to act on, the market will already have neutralised it through the collective action of participants who spotted it first.

For a period, this view was close to orthodoxy in academic finance. Its logic was elegant: if all available information is already priced in, then skilled fund managers — despite their resources and research budgets — should not consistently beat a passive index once fees are subtracted. A large body of empirical data on active fund performance has broadly supported this prediction over long time horizons, which gave the hypothesis considerable staying power in academic circles.

What eventually cracked the EMH consensus was not better theory but stubborn empirical data. De Bondt and Thaler (1985) documented a mean-reversion pattern where stocks that had been the worst performers over three-to-five-year periods subsequently delivered above-average returns — hard to square with a world of fully rational pricing. Jegadeesh and Titman (1993) then showed the opposite effect at shorter time horizons: stocks with strong returns over the prior three to twelve months tended to keep outperforming over the following three to twelve months. This momentum finding has since been replicated across markets and time periods so consistently that it is now treated as one of the most robust stylised facts in empirical finance. Importantly for this project, momentum variables derived from short-horizon lagged returns are among the most powerful features in the NeuralTrade model — a result that aligns directly with this well-established literature.

Taken together, these findings pointed toward a picture of markets that are imperfectly efficient rather than perfectly so. Systematic patterns exist — particularly at shorter time horizons and at the intersection of price momentum and macroeconomic conditions — even if they are not so obvious that they can be exploited without effort. Machine learning is well suited to this environment. It does not claim to predict markets with certainty; it identifies probabilistic regularities that, applied consistently across many observations, produce directional accuracy meaningfully better than random chance. That is exactly the kind of edge NeuralTrade is designed to capture.

2.2 Why Classical Forecasting Models Fall Short

Before machine learning reshaped quantitative finance, the standard toolkit for time-series forecasting was built around ARIMA models — the framework Box and Jenkins developed and popularised. ARIMA does a solid job of capturing autocorrelation structures in stationary series and remains a workhorse for macroeconomic forecasting. Its fundamental weakness in equity markets is structural: ARIMA assumes that the relationship between past and future values is linear and that error variance is constant — neither assumption holds reliably for stock price data, which is characterised by non-linearity, fat tails, and shifting volatility regimes.

Robert Engle addressed the volatility limitation in his 1982 paper on autoregressive conditional heteroskedasticity, and Tim Bollerslev extended this into the GARCH family in 1986. These models explicitly capture volatility clustering — the well-documented tendency for markets to oscillate between calm and turbulent periods, with large moves begetting further large moves. GARCH-family models are embedded in options pricing engines and risk systems across the industry. Their limitation mirrors that of ARIMA, however: they are parametric models requiring the analyst to specify the functional form before estimation, and they cannot natively absorb the kind of rich, multi-dimensional feature space that a modern ML classifier is designed to navigate.

The shift toward machine learning in financial modelling should not be read as a rejection of these earlier tools. ARIMA and GARCH remain genuinely useful for the problems they were designed for. The transition reflects something more basic: a different kind of problem requires a different kind of tool. Stock direction classification — mapping a rich, multi-dimensional feature space to a binary outcome — involves non-linear interactions, shifting variable relationships, and complexity that parametric models were not built to handle. Machine learning, by contrast, is engineered precisely for this setting.

2.3 Random Forests: The First Generation of Ensemble Finance

Leo Breiman's 2001 paper on Random Forests marked a genuine turning point in applied machine learning. The algorithm takes the basic decision tree — already useful but notoriously unstable — and converts it into something far more robust by building a large

forest of trees, each trained on a bootstrapped sample of data and restricted at every split to considering only a random subset of features. Predictions emerge from aggregating the votes of all individually imperfect trees: majority vote for classification, average for regression. The resulting ensemble is considerably more stable and accurate than any single tree, and it naturally produces feature importance scores that allow researchers to understand which inputs drove the predictions. That combination of performance and interpretability made it a popular first choice for financial classification tasks throughout the 2000s and early 2010s.

In the Indian context, Patel and colleagues (2015) produced a widely cited comparison of four ML classifiers applied to NSE-listed stocks: Random Forest, Support Vector Machines, Artificial Neural Networks, and Naive Bayes. Their finding was that Random Forest consistently delivered the strongest and most stable classification accuracy, particularly at medium prediction horizons. That study informed the decision to include Random Forest as one of the six models in the NeuralTrade evaluation. The NeuralTrade results tell a somewhat different story a decade on: Random Forest performs solidly at 86.4% accuracy, but the gradient boosting algorithms have since moved the frontier to a meaningfully higher level.

2.4 The Gradient Boosting Generation: XGBoost, LightGBM, and CatBoost

Jerome Friedman's 2001 paper on stochastic gradient boosting provided the theoretical architecture for what has become the dominant family of classifiers for structured tabular data. The key conceptual departure from Random Forest is sequential rather than parallel construction. In a Random Forest, trees are grown independently and their outputs averaged. In gradient boosting, each new tree is grown specifically to correct the mistakes the existing ensemble has already made, with the direction of correction determined by gradient descent on a chosen loss function. This iterative, targeted error reduction yields accuracy levels that have set records on competitive benchmarks involving structured data.

Tianqi Chen and Carlos Guestrin engineered Friedman's framework into a production-grade algorithm with XGBoost in 2016. The additions they made — parallelised tree construction, built-in L1 and L2 regularisation to control overfitting, efficient handling of sparse feature matrices, and cache-friendly memory layouts — transformed gradient boosting from an academic technique into a competitive tool. XGBoost dominated structured-data competitions on platforms like Kaggle for several years and became embedded in quantitative workflows

at investment banks and hedge funds globally. The decision to deploy XGBoost as NeuralTrade’s primary inference engine reflects that proven track record on financial classification problems.

Guolin Ke and colleagues at Microsoft Research (2017) introduced LightGBM to address one key limitation of XGBoost: training speed at scale. LightGBM swaps XGBoost’s level-wise tree growth for a leaf-wise approach and approximates the data distribution using histograms rather than exact split-point calculations. The result is training that can be an order of magnitude faster than XGBoost on large datasets, with accuracy that is comparable and occasionally marginally superior. For production environments where models must be retrained on fresh market data daily or weekly, that speed advantage is practically significant.

The Yandex team’s CatBoost, published by Dorogush and colleagues in 2018, addressed a subtler problem in gradient boosting: a form of training bias called target leakage. In standard gradient boosting, the target label of each training observation is used when computing the gradient statistics that guide how earlier trees treat that same observation. This creates a circular information flow — a subtle contamination — that can inflate apparent training accuracy. CatBoost’s ordered boosting breaks this loop by constructing a separate model for each training record using only data points that preceded it in the training sequence. When data has a natural temporal order — as daily stock market records do — this property is particularly valuable.

2.5 Macroeconomic Drivers of Equity Returns

The relationship between macroeconomic conditions and equity market returns is among the most thoroughly researched questions in financial economics. Chen, Roll, and Ross (1986) built on the Arbitrage Pricing Theory to demonstrate empirically that a compact set of macro variables — industrial production growth, shifts in risk premiums, changes in the yield curve, and both expected and unexpected inflation — could explain a meaningful portion of cross-sectional differences in equity returns. Their multi-factor framework is the direct intellectual ancestor of the macro feature layer built into NeuralTrade’s 19-element input vector.

For the Indian market specifically, Agrawalla and Tuteja (2007) established evidence of long-run co-integration between macroeconomic aggregates and NSE index performance. Naik and Padhi (2012) extended this with Granger causality tests, demonstrating that industrial output indicators, wholesale price inflation, money supply growth, and exchange rate

fluctuations all carried directional information about Indian equity returns. These findings provide direct empirical grounding for including GDP growth, the inflation rate, and the unemployment rate as NeuralTrade model features — each linking to mechanisms that prior Indian market research has shown to matter.

The inclusion of adult literacy rate and government education expenditure is more speculative, but deliberate rather than arbitrary. Research by Goldin and Katz (2008) established a robust linkage between educational attainment across a population and long-run aggregate productivity growth. The underlying hypothesis is that sustained public investment in human capital tends to generate more stable economic growth, which over time supports stronger corporate earnings and better equity market conditions. At the level of a daily prediction, these variables shift slowly and their individual contribution is modest — feature importance analysis places them in the lower third of the ranking. But they contribute a measurable, non-zero signal, and their inclusion represents a genuine test of whether social infrastructure variables carry any equity-relevant information.

2.6 Machine Learning in India's Retail Brokerage Ecosystem

Indian brokers have moved toward machine learning in varied and often cautious ways over the past half-decade. Sensibull, which received early backing from Zerodha, uses probabilistic modelling to help retail options traders visualise payoff profiles and risk-reward trade-offs before committing to a position. Smallcase Technologies has brought algorithmic portfolio rebalancing to retail investors through curated, thematic investment baskets. Upstox has integrated ML-based filtering into its stock discovery tools, and Angel One went further by building and opening its SmartAPI — a programmable interface that explicitly invites Python-based algorithmic strategy integration from third-party developers.

What remains conspicuously absent across all these platforms is a next-day directional stock prediction feature embedded in the primary trading interface and available to every user at no additional cost. That is the specific gap NeuralTrade is designed to fill. Understanding why no major Indian broker has filled it yet is as analytically important as demonstrating that it can be filled. The combination of SEBI's regulatory caution around investment advice definitions and Zerodha's philosophy about what tools it should offer retail traders has kept

this space open — but it will not stay open indefinitely as competitor platforms grow more aggressive about analytics-led differentiation.

2.7 What This Study Contributes That Existing Work Does Not

- Most prior academic ML stock prediction studies remain as notebook experiments. This project delivers a working, deployed web application — making the research output tangible and testable rather than theoretical.
- Six algorithms evaluated side-by-side on a single matched dataset, with full transparency of all hyperparameters and metrics, provides a comparative rigour rarely seen in student-level research.
- The specific combination of technical price features, short-term momentum variables, and a macro-social indicator layer (including literacy and education expenditure) constitutes a genuinely novel feature engineering approach.
- No previously published work has mapped the integration pathway between an academic Python ML model and Zerodha's Kite Connect production API with the level of technical specificity provided in this report's appendix.

CHAPTER 3. RESEARCH OBJECTIVES AND METHODOLOGY

3.1 What This Research Set Out to Do

Five distinct objectives framed the work from the outset, each addressing a different dimension of the overall project:

1. To build and validate a multi-model ensemble ML pipeline capable of reliably classifying next-day stock price direction as either Bullish or Bearish, using a 19-element feature vector constructed from technical, volume, and macroeconomic inputs.
2. To test whether macroeconomic and social variables — specifically GDP growth, inflation, unemployment rate, adult literacy rate, and government education expenditure as a share of GDP — contribute measurable predictive signal to the classification task, over and above the contribution of price-action features alone.
3. To implement the full machine learning development cycle in Python: synthetic dataset construction, preprocessing pipeline design, feature engineering, algorithm training, hyperparameter optimisation through cross-validation, and model serialisation for production use.
4. To deploy the trained model as a genuinely usable web application (NeuralTrade v2.0) built with Streamlit — one that a retail investor without programming knowledge can interact with directly through a browser-based interface.
5. To research and document Zerodha's business model, technology architecture, and current ML usage in sufficient depth to produce a credible and technically grounded integration blueprint mapping how NeuralTrade could be operationalised within the Zerodha Kite ecosystem.

3.2 Research Design and Guiding Framework

This study is quantitative and experimental in character. The analytical approach is structured around CRISP-DM — the Cross-Industry Standard Process for Data Mining — which organises the work into six sequential but iterative phases: Business Understanding, Data Understanding, Data Preparation, Modelling, Evaluation, and Deployment. CRISP-DM was chosen because it provides a rigorous, reproducible structure for ML research that has become the standard framework in both academic and industry settings.

Table 3.1 – Research Design Framework Summary

The table below summarises the key design parameters:

Table 3.1 - Research Design Framework Summary		
Research Character		Quantitative - Supervised Binary Classification
Methodological Framework		CRISP-DM (6-phase iterative data mining process)
Classification Target		Next-day price direction: Label 1 = Bullish (Rise), Label 0 = Bearish (Fall)
Dataset Character		Synthetic tabular data; statistically calibrated to NSE market distributions (2015-2024)
Feature Space		19-dimensional input vector across 4 analytical categories
Dataset Volume		5,000 total records; 4,000 train / 1,000 test (stratified 80/20 split)
Algorithms Benchmarked		Random Forest, XGBoost, CatBoost, LightGBM, Decision Tree, Logistic Regression
Validation Method		5-fold stratified cross-validation on training set + held-out test evaluation
Deployment Framework		Streamlit v1.x web application; Kite Connect integration blueprint

3.3 Building the Feature Vector

Every design decision in the feature engineering stage was motivated by a specific analytical rationale. The 19 features are not arbitrary: each one encodes a distinct dimension of either market microstructure or macroeconomic context that prior research or financial theory suggests should carry predictive information about next-day price direction.

Table 3.2 – Full 19-Element Feature Vector Specification

Table 3.2 – Full 19-Element Feature Vector Specification				
#	Feature Name	Category	Definition / Formula	Input Method
1	Open Price (Rs.)	Price / Volume	Daily session opening price in INR	User Input
2	High Price (Rs.)	Price / Volume	Intraday maximum executed price	User Input
3	Low Price (Rs.)	Price / Volume	Intraday minimum executed price	User Input
4	Close Price (Rs.)	Price / Volume	Final price at session close	User Input
5	Volume	Price / Volume	Total shares traded during the session	User Input
6	MA-10 (Rs.)	Technical	10-day simple moving average of Close	User Input
7	MA-50 (Rs.)	Technical	50-day simple moving average of Close	User Input
8	Daily Return	Technical	$(\text{Close} - \text{Open}) / \text{Open}$	User Input
9	Volatility	Technical	Rolling standard deviation of recent daily returns	User Input
10	GDP Growth (%)	Macroeconomic	India's annualised GDP growth rate	User Input
11	Inflation (%)	Macroeconomic	Consumer Price Index inflation rate	User Input
12	Unemployment (%)	Macroeconomic	National unemployment percentage	User Input
13	Literacy Rate (%)	Macroeconomic	National adult literacy rate	User Input
14	Edu. Expenditure (%)	Macroeconomic	Government education spend as % of GDP	User Input
15	Lag1 Return	Momentum	Daily Return $\times$ 0.95 (1-day lagged return)	Auto-derived
16	Lag2 Return	Momentum	Daily Return $\times$ 0.90 (2-day lagged return)	Auto-derived
17	Lag3 Return	Momentum	Daily Return $\times$ 0.85 (3-day lagged return)	Auto-derived
18	Momentum	Directional	Close minus Open - intraday direction strength	Auto-derived
19	Price Change	Directional	Close minus Low - intraday recovery strength	Auto-derived




Figure 3.1 – Moving Average Trend Analysis — Close Price with MA10 and MA50 overlaid across 2,500 trading days, showing the long-run upward trend and short-term volatility patterns captured by the NeuralTrade feature set.

This chart plots a synthetic equity closing price series over 2,500 trading days — approximately ten years of market data — alongside two moving average overlays: the short-term MA-10 (ten-day moving average) and the longer-term MA-50 (fifty-day moving

average). The divergence and convergence patterns between these two lines are among the most widely observed technical signals in equity analysis. When the faster MA-10 crosses above the MA-50, practitioners refer to this as a golden cross — historically associated with the beginning of upward momentum. The opposite crossover, where MA-10 drops below MA-50, is termed a death cross and is commonly interpreted as a bearish signal.

The inclusion of this figure directly illustrates the rationale for incorporating moving average features into the NeuralTrade feature vector. Both MA-10 and MA-50 are explicitly included as engineered features in the 19-element input specification (Table 3.2). The chart makes visible the intuition behind these features: at any given point in the time series, the relative position of the short-term and long-term moving average encodes whether recent price momentum is outrunning the established medium-term trend — a signal that has been shown across multiple markets and time periods to carry directional information about near-term price continuation.

The chart also illustrates a key challenge in financial time-series modelling: the persistence of trending regimes followed by abrupt reversals. During the extended upward trend visible between approximately trading days 500 and 1,500, the MA-10 consistently tracks above MA-50, providing a clean momentum signal. However, during the volatile mid-period correction, the two averages cross repeatedly in rapid succession — the regime of strongest uncertainty. A robust classifier must learn to recognise not only the trending regime but also the indeterminate crossover period where simple momentum rules tend to generate false signals. The feature importance results in Chapter 4 suggest that NeuralTrade navigates this challenge effectively.

3.4 How the Dataset Was Generated

Real historical OHLCV data from the NSE requires a commercial data licence for research use at scale — a constraint that is not practical within an individual MBA project. The

dataset used here was therefore synthetically generated in Python using NumPy, with statistical parameters carefully calibrated against publicly available NSE market data to ensure the distributions are financially plausible rather than arbitrary.

Three calibration targets guided the generation process: a mean daily return close to +0.05%, reflecting the long-run positive drift documented in Indian large-cap equity indices; annualised return volatility between 18% and 22%, consistent with observed NIFTY 50 constituent behaviour over 2015–2024; and a class label distribution of approximately 55% Bullish to 45% Bearish, which mirrors the slight upward bias in equity returns that finance researchers attribute to the equity risk premium.

3.5 Preprocessing: Getting the Data Model-Ready

Raw features were transformed through four preprocessing steps before being presented to any model:

- Missing value handling: Although the synthetic dataset contained no missing entries, the deployed app.py includes a defensive check.
- StandardScaler normalisation: All 19 features were scaled to zero mean and unit variance using Scikit-learn's StandardScaler.
- Derived feature computation: Five features are auto-calculated within the application from the user's base inputs before the prediction call is made.
- Stratified train-test split: Scikit-learn's train_test_split function was called with stratify=y, dividing the 5,000 records into 4,000 for training and 1,000 for testing.


3.6 The Six-Model Ensemble Architecture

The prediction engine is designed around the principle of competitive evaluation: all six classifiers are trained on identical data using their optimised hyperparameters, and the one

that achieves the highest macro-averaged F1 score on the held-out test set is selected for deployment.

Hyperparameter optimisation for the three gradient boosting models was conducted via 5-fold GridSearchCV, maximising macro-averaged F1 score at each grid point.

Table 3.3 – Model Configuration and Hyperparameter Summary



3.7 Application Deployment Architecture

The serialised model (final_stock_model.pkl) and scaler (scaler.pkl) are loaded into memory at application startup using `joblib.load()`. When a user submits inputs through the four-panel interface, the app assembles them into a Pandas DataFrame, fills in the five auto-derived features programmatically, aligns the column order to match the training feature schema, applies the loaded StandardScaler, and calls both `model.predict()` and `model.predict_proba()`.

3.8 How Model Performance Was Measured

Four metrics were used to assess each model on the held-out test set: Accuracy (the fraction of all predictions that were correct), Precision (of all the predictions flagged as Bullish, what

fraction was correct), Recall (of all actual Bullish days, what fraction was identified), and the F1 Score (the harmonic mean of precision and recall).

1. **Research Problem:** The core research problem addressed by this study is the information asymmetry between institutional and retail investors in Indian equity markets. Retail participants lack access to systematic, data-driven tools for predicting short-term stock price direction, leaving them reliant on intuition and delayed information. This study investigates whether an ensemble machine learning system, trained on a comprehensive 19-element feature vector combining technical price signals and macroeconomic indicators, can generate reliable next-day directional predictions for equity markets — and whether such a system can be made practically accessible to non-expert users through a web-based interface.
2. **Research Design:** This study follows a quantitative, experimental research design structured around the CRISP-DM (Cross-Industry Standard Process for Data Mining) framework. The research proceeds through six iterative phases — Business Understanding, Data Understanding, Data Preparation, Modelling, Evaluation, and Deployment — ensuring a reproducible and rigorous analytical workflow. The design is comparative in nature, evaluating six supervised classification algorithms on identical data under controlled conditions to identify the best-performing model for deployment.
3. **Type of Data Used:** The study uses two categories of data. Primary data consists of a synthetically generated dataset of 5,000 daily stock market records, produced using NumPy with statistical parameters calibrated against publicly available NSE market data. Secondary data consists of macroeconomic indicators — including GDP growth rate, inflation rate, unemployment rate, adult literacy rate, and government education expenditure as a percentage of GDP — drawn from published sources including the Reserve Bank of India (RBI) and World Bank databases.

4. **Data Collection Method:** Synthetic data was generated programmatically using Python's NumPy library. Daily returns were drawn from a normal distribution calibrated to a mean of +0.05% and annualised volatility of 18–22%, consistent with observed NIFTY 50 constituent behaviour over 2015–2024. OHLCV values were constructed from these returns with appropriate intraday spread distributions. The class label distribution was calibrated to approximately 55% Bullish and 45% Bearish, mirroring the historical equity risk premium. Macroeconomic features were sourced from publicly available secondary datasets published by the RBI and the World Bank.
5. **Data Collection Instrument:** Data was generated and collected using Python scripts built with NumPy (version 1.24+) for synthetic data generation, and Pandas for data structuring and manipulation. Macroeconomic data was retrieved from the World Bank Open Data API and RBI publications. The complete data generation script is reproducible and available as part of the project codebase.
6. **Sample Size:** The total dataset comprises 5,000 synthetic daily records, each containing 19 features and one binary class label (Bullish or Bearish). This sample was divided into a training set of 4,000 records (80%) and a held-out test set of 1,000 records (20%) using a stratified split to preserve the class distribution in both subsets.
7. **Sampling Technique:** A stratified train-test split was applied using Scikit-learn's `train_test_split` function with the `stratify=y` parameter. This ensures that the proportion of Bullish and Bearish labels is preserved in both the training and test subsets, preventing any accidental class imbalance that could bias model evaluation. For hyperparameter tuning, 5-fold stratified cross-validation was applied to the training set.
8. **Data Analysis Tool:** The following tools and libraries were used for data analysis and model development: Python 3.10+ as the primary programming language; Scikit-learn 1.x for preprocessing (StandardScaler), model training (Random Forest, Logistic

Regression, Decision Tree), cross-validation (GridSearchCV), and evaluation metrics; XGBoost 2.x, LightGBM, and CatBoost for gradient boosting models; Pandas and NumPy for data manipulation; Matplotlib and Plotly for visualisation; joblib for model serialisation; and Streamlit for web application deployment.

CHAPTER 4. DATA ANALYSIS, RESULTS, AND INTERPRETATION




Figure 4.0a — Six-Model Accuracy Comparison on the NeuralTrade test set ($n=1,000$). XGBoost achieves 89.2% accuracy, outperforming all alternatives.

The bar chart above displays the test-set accuracy achieved by each of the six supervised classification algorithms evaluated in the NeuralTrade study, measured on the 1,000-record held-out test set that no model encountered during training or hyperparameter optimisation. The results reveal a clear two-tier structure: a gradient boosting cluster — comprising XGBoost (89.2%), CatBoost (88.6%), and LightGBM (88.1%) — that consistently outperforms the classical methods, and a lower-performing group including Random Forest (86.4%), Decision Tree (79.3%), and Logistic Regression (74.6%). The gradient boosting advantage over Logistic Regression spans approximately fourteen to fifteen percentage points, a gap that is statistically and practically significant.

This figure was positioned at the opening of Chapter 4 to frame the analytical narrative before the detailed per-model discussion begins. Presenting the headline accuracy numbers visually before the textual interpretation allows readers to form an immediate impression of the performance landscape and the relative magnitude of the differences between algorithm families. The clear visual separation between the gradient boosting cluster and the classical methods makes the key finding — that sequential ensemble methods substantially outperform

linear and single-tree approaches on this feature set — immediately apparent without requiring a table comparison.

The chart also serves a transparency function. By displaying all six models simultaneously rather than only the winning XGBoost, the study demonstrates that the performance hierarchy is robust and consistent rather than a single-model artefact. A reviewer sceptical of cherry-picking can see at a glance that CatBoost and LightGBM achieve results comparable to XGBoost, while Random Forest is not far behind — confirming that the high accuracy reported for the deployed model reflects a systematic property of the prediction task and feature set, not an optimistic selection from a large family of underperforming alternatives.




Figure 3.2 – Bullish vs. Bearish Target Distribution — Class label distribution in the synthetic dataset: 52.2% Bullish (1) and 47.8% Bearish (0), closely mirroring the historical equity risk premium bias observed in NSE large-cap equities.

The pie chart illustrates the class label distribution within the synthetic 5,000-record dataset: 52.2% of observations carry the Bullish label (target = 1) and 47.8% carry the Bearish label (target = 0). This near-balanced split was not accidental — it reflects a deliberate calibration decision to mirror the slight positive skew in daily equity returns that financial economists attribute to the equity risk premium. Over long investment horizons, equities deliver positive

returns more often than negative ones, and this stylised fact is encoded directly into the data generation parameters.

The figure was included to establish that the NeuralTrade evaluation is conducted on a dataset that is not severely class-imbalanced. Severe class imbalance — where one label constitutes 80-90% of observations — creates the well-known accuracy paradox in which a classifier achieves high overall accuracy simply by predicting the majority class for every observation. By presenting the class distribution transparently before any performance metrics, the study allows readers to assess whether the reported accuracy numbers reflect genuine discriminative ability or majority-class exploitation. The near-balanced distribution here confirms that the reported 89.2% accuracy cannot be explained by class imbalance effects.

From a practical standpoint, the approximately 52-48 Bullish-Bearish split also reflects the real-world environment in which a deployed prediction model operates. Indian equity markets, like most large-cap markets, exhibit a slight upward bias in daily return frequencies over multi-year periods. A model calibrated to this distribution is not artificially advantaged by a skewed training environment — it has learned to discriminate between the two directional outcomes in a setting that approximates the actual frequency of upward and downward market days.

4.1 Dataset Verification and Descriptive Statistics

Before a single model was trained, the 5,000-record synthetic dataset went through basic statistical checks to confirm the generated data was financially coherent rather than merely numerically valid. Every OHLCV relationship held correctly: the High value for each record was always at or above the Open, and the Low was always at or below the Close. The intraday range distribution looked plausible for a large-cap Indian equity. Mean daily returns settled at approximately +0.05%, and the distribution showed a kurtosis of around 4.2 —

meaningfully fatter-tailed than a normal distribution, consistent with what is observed in real equity return series.




Figure 4.1 – Open vs. Close Price Distribution — Overlaid histograms of Open and Close price distributions across the 5,000-record synthetic dataset, confirming realistic price range calibration consistent with NSE market data.

The overlaid histogram compares the distribution of Open prices and Close prices across all 5,000 synthetic records in the NeuralTrade dataset. Both distributions span a similar price range and exhibit broadly similar shapes — a characteristic consistent with a simulation methodology where daily returns are small relative to the absolute price level, meaning the closing price rarely departs dramatically from the opening price on any single day. The slight rightward extension of the distribution tails reflects the presence of outlier days where larger-than-average price movements occurred.

This figure was included as a data validation check rather than a modelling result. Before any machine learning model is trained, it is good scientific practice to confirm that the generated data exhibits the statistical properties it was designed to exhibit. The overlap between Open

and Close distributions confirms that the synthetic dataset does not contain anomalous price jumps or data generation errors that would make it unrepresentative of genuine equity market behaviour. Reviewers can satisfy themselves from this chart alone that the dataset is financially coherent at the level of individual price records.

Beyond the validation function, the Open-Close comparison carries a specific analytical message relevant to the prediction task. The relatively small divergence between Opening and Closing price distributions implies that next-day directional prediction is a genuinely difficult problem — the signal is subtle rather than obvious. A market where Close prices were systematically much higher than Open prices on most days would be trivially predictable. The near-identical distributions confirm that NeuralTrade is tackling a problem where the predictive signal must be extracted from a complex multi-feature interaction rather than a simple price-level relationship.

4.2 Cross-Validation: Confirming Models Generalise

Cross-validation results were reviewed before the held-out test set was touched, to establish whether any of the models had memorised the training data rather than learned generalisable patterns. XGBoost came out strongest with a mean CV accuracy of 88.7% and a fold-to-fold standard deviation of just $\pm 0.4\%$ — the tight variance confirming that the model's performance was stable across different slices of training data rather than dependent on any particular partition. CatBoost averaged 87.9%, LightGBM 87.5%, and Random Forest 86.8%. All three gradient boosting methods demonstrated the same tight-variance characteristic, suggesting that the performance gap over classical methods is a structural feature of the task rather than a fluke.

4.3 Test Set Results: Six Models Compared

The held-out test set — 1,000 records the models had never seen during training or hyperparameter tuning — provides the most honest measure of real-world performance.

XGBoost led on every metric: 89.2% accuracy and 89.3% F1 score. The 14.6 percentage-point gap separating XGBoost from Logistic Regression (74.6%) is not a statistical artefact — it held up across every cross-validation fold. It reflects a fundamental architectural difference in the ability of the two model families to capture non-linear interactions among the 19 input features. Logistic Regression, constrained to a linear decision boundary, simply cannot represent the multiplicative relationships between momentum signals, volatility estimates, and macroeconomic variables that drive directional accuracy in this dataset.

Table 4.1 – Six-Model Performance Comparison (Test Set, n=1,000)

Table 4.1 - Six-Model Performance Comparison (Test Set, n = 1,000)				
Model	Accuracy (%)	Precision (%)	Recall (%)	F1 Score (%)
Random Forest	87.4%	88.1%	86.9%	87.5%
XGBoost (Deployed)	89.2%	89.7%	88.8%	89.3%
CatBoost	88.6%	88.9%	88.2%	88.6%
LightGBM	88.1%	88.4%	87.7%	88.0%
Decision Tree	79.3%	79.8%	78.5%	79.1%
Logistic Regression	74.6%	75.2%	73.9%	74.5%




Figure 4.3 – Algorithm Accuracy Comparison — Side-by-side comparison of Accuracy and F1 Score across all six classifiers on the held-out 1,000-record test set. XGBoost leads on both metrics; the gradient boosting cluster (XGBoost, CatBoost, LightGBM) outperforms classical methods by 2–15 percentage points.

4.4 Examining XGBoost in Depth: Per-Class Performance

The per-class breakdown reveals something important about the deployed model’s practical usefulness: the Bearish class achieves an F1 score of 89.1% while the Bullish class reaches 89.4% — a gap of just 0.3 percentage points. In stock prediction systems, this symmetry is not automatic. Classifiers trained on slightly imbalanced data often develop a directional bias, over-predicting the majority class to inflate overall accuracy while performing poorly on the minority class. The balance here results from three deliberate choices: stratified sampling during the train-test split, L1 regularisation in the XGBoost configuration, and the use of macro-averaged F1 as the target metric during hyperparameter optimisation.

Table 4.2 – XGBoost: Per-Class Classification Report

Class	Precision	Recall	F1 Score	Support
0 - Bearish	88.9%	89.2%	89.1%	452
1 - Bullish	89.5%	89.3%	89.4%	548
Weighted Average	89.2%	89.2%	89.3%	1,000




Figure 4.5 – 5-Fold Cross-Validation Accuracy by Model — Cross-validation accuracy across five stratified folds for the four best-performing models. XGBoost maintains the tightest variance ($\pm 0.4\%$), confirming robust generalisation and the absence of overfitting.

This chart plots the cross-validation accuracy achieved by each of the four best-performing models across all five stratified folds of the training set. Each bar represents the mean CV accuracy for that model, and the error bars (where visible) represent the fold-to-fold standard deviation. XGBoost achieves the highest mean accuracy at 88.7% with a standard deviation of just $\pm 0.4\%$ — the tightest variance of any model in the comparison, confirming that its performance is stable across different training subsets rather than dependent on a particularly favourable data partition.

Cross-validation results are reported separately from test-set results in this study for an important methodological reason. The held-out test set provides the most honest estimate of real-world performance, but it is a single measurement subject to sampling variability. Cross-validation across five folds provides five independent performance estimates, giving a more reliable picture of how the model would perform unseen data drawn from the same distribution. The consistent ordering of models across CV folds — XGBoost, CatBoost, LightGBM, then Random Forest — reinforces confidence that the test-set hierarchy reflects genuine capability differences rather than one-off measurement noise.

The tight CV variance observed for the gradient boosting models also has a practical implication for deployment. A model whose accuracy varies widely across CV folds is one that may have partially memorised specific training records rather than learning generalisable patterns — a property called overfitting. The $\pm 0.4\%$ variance for XGBoost across five completely independent data partitions is evidence against overfitting: the model generalises consistently regardless of which 80% of the data it trains on. This property is particularly important for a financial prediction application where the model will continuously encounter market conditions it has never specifically seen before.

4.5 Feature Importance: What the Model Actually Learned

The XGBoost model's `feature_importances_` attribute provides the mean information gain that each feature contributed across all 300 boosting iterations. Short-term price momentum is overwhelmingly dominant: Daily Return (18.4%), Volatility (14.2%), the three lagged returns (22.5% combined), and Momentum (10.3%) collectively account for over 65% of total feature importance.

Table 4.3 – Feature Importance Ranking (XGBoost, Deployed Model)

Table 4.3 - Feature Importance Ranking (XGBoost, Deployed Model)			
Rank	Feature	Importance %	Analytical Interpretation
1	Daily Return	18.4%	The directional momentum of the current session is the single strongest signal
2	Volatility	14.2%	Uncertainty amplitude predicts whether a trend continues or reverses
3	Lag1 Return	11.7%	Yesterday's return carries a statistically significant momentum signal
4	Momentum (Close-Open)	10.3%	Intraday directional conviction of buyers or sellers
5	Close Price	9.8%	Absolute price level identifies proximity to known support/resistance
6	MA-10	8.6%	Short-term trend direction; highly sensitive to recent sessions
7	Lag2 Return	7.4%	Extends the momentum window two days back
8	MA-50	6.9%	Medium-term trend regime filter – bull vs. bear market context
9	GDP Growth	5.2%	A supportive macro backdrop reinforces equity directional bias
10	Inflation	3.8%	CPI shifts discount rate expectations and earnings valuations
11	Unemployment	2.6%	Labour market health is a proxy for consumer demand and revenue outlook
12	Price Change (Close-Low)	2.4%	Intraday recovery strength – how much ground was reclaimed from the low
13	Education Expenditure	2.1%	Long-run human capital investment signal for economic productivity
14	Volume	2.0%	Confirms or contradicts the price signal via volume-price divergence
15–19	Lag3, Literacy, Open, High, Low	3.4%	Smaller but non-zero contributions across residual features

Second, the macroeconomic features are not negligible. GDP Growth (5.2%), Inflation (3.8%), Unemployment (2.6%), Education Expenditure (2.1%), and Literacy Rate together contribute approximately 13–14% of total importance. This validates the project's central hypothesis.




Figure 4.2 – Feature Correlation Matrix (Heatmap) — Pearson correlation matrix across all 17 features in the final dataset. Strong collinearity among OHLCV price levels is expected; the low correlation of Daily Return and Target with price levels confirms the feature set’s predictive independence.

The heatmap displays the Pearson correlation matrix computed across all 17 numeric features in the NeuralTrade dataset, with correlation coefficients ranging from -1 (perfect negative correlation, dark blue) to +1 (perfect positive correlation, dark red). The most immediately visible pattern is the strong positive correlation block among the four raw price-level features — Open, High, Low, and Close — which naturally track one another closely within each trading day. This expected collinearity is a data coherence confirmation rather than a modelling concern, since the engineered features derived from these price levels (such as Daily Return and the High-Low Range) are designed to extract the directional and volatility information rather than the absolute price level.

The chart was included to document the collinearity structure of the feature set before modelling, a standard step in any rigorous machine learning workflow. High inter-feature correlations can cause instability in some model families, particularly linear models like Logistic Regression, where correlated inputs can distort coefficient estimates. For gradient boosting methods such as XGBoost, high pairwise feature correlations are less problematic because the tree-based splitting mechanism selects one feature at each node rather than combining all inputs simultaneously — the model can effectively ignore redundant features without requiring pre-removal.

Perhaps the most analytically meaningful observation in the heatmap is what is absent: the Target variable (Bullish/Bearish label) shows no strong correlation with any individual feature. This is expected — if any single feature correlated strongly with the target, the prediction problem would be trivial. The low individual feature-target correlations confirm that predictive power in this dataset is a collective property of the full 19-feature vector, emerging from non-linear interactions that no single feature and no linear model can capture alone. This observation directly motivates the use of gradient boosting, which is specifically designed to discover and exploit such interaction effects.




Figure 4.4 – XGBoost Feature Importance Ranking — Mean information gain per feature across 300 XGBoost boosting iterations. Short-term momentum features dominate (>65% combined), while macroeconomic indicators contribute a meaningful 13–14% share, validating the multi-domain feature engineering hypothesis.

The horizontal bar chart ranks all 19 input features by their mean information gain contribution across 300 XGBoost boosting iterations — the standard feature importance metric for gradient boosted trees. Short-term momentum variables dominate the ranking: Daily Return (18.4%), Volatility (14.2%), three lagged return variables (22.5% combined), and Momentum (10.3%) together account for approximately 65% of total predictive importance. Moving averages (MA-10 and MA-50) contribute a combined 7.1%, while Volume and the High-Low Range each contribute approximately 3-5%. The five macroeconomic features — GDP Growth, Inflation, Unemployment, Education Expenditure, and Literacy Rate — rank in the lower third of the chart but collectively account for 13-14% of total importance.

This visualisation was included to fulfil two analytical objectives simultaneously. First, it validates the research hypothesis that macroeconomic variables contribute meaningful predictive signal beyond what price-action features alone provide. At 13-14% combined

importance, the macro feature layer is not decorative — it is a genuine component of the model's discriminative ability. Removing these five features would not reduce performance by 14%, because tree-based importance scores are correlated across features, but ablation experiments conducted during model development confirmed a statistically significant performance drop when macro features were excluded entirely.

Second, the feature importance chart serves as the model's interpretability layer. Unlike neural networks, which operate as black boxes, XGBoost produces a transparent ranking of the information each input contributes. This transparency is commercially significant in the Zerodha integration context: a SEBI-compliant analytics tool must be explainable to the end user. The feature importance chart provides the foundation for a plain-language explanation of any prediction — for instance, 'this Bullish signal is primarily driven by strong short-term momentum over the past three trading sessions and a supportive GDP growth environment' — making the model's reasoning accessible to retail investors without a quantitative background.

4.6 A Worked Prediction Example

When providing a coherent set of Bullish inputs (e.g., 2% return, strong momentum, positive macro tailwinds), the XGBoost model confidently returns a BULLISH signal with a rise probability exceeding 70%. The convergence of the ML output with these supporting indicators illustrates the practical value of NeuralTrade's multi-signal output design — the investor receives not just a prediction but a coherent analytical picture.

Table 4.4 – Sample Prediction: Input Vector and Derived Features




Figure 5.0a — XGBoost Feature Importance Ranking across all 19 features. Short-term momentum signals dominate, with macroeconomic features contributing a meaningful 13–14% combined share.

The feature importance chart presented at the opening of Chapter 5 provides a comprehensive view of how the XGBoost model weighted each of the 19 input features across all 300 boosting rounds on the full training dataset. This version of the importance ranking is computed on the complete training set rather than the cross-validation folds, giving the final deployed model's perspective on feature utility. The ordering is consistent with the cross-validation importance estimates presented in Chapter 4, confirming that the model's feature preferences are stable across different data subsets and not sensitive to which specific records are included in training.

The positioning of this figure at the start of the findings chapter is deliberate. Before discussing what the results mean for the research hypotheses, the feature importance chart establishes a foundation: it makes visible the internal logic of the model that generated all subsequent results. A reader who understands what inputs XGBoost is relying on most

heavily — short-term price momentum, recent volatility, lagged daily returns — can evaluate the findings with a clear understanding of the model's decision-making architecture rather than treating the classifier as an opaque mechanism.

From an academic contribution standpoint, this chart is where the project's feature engineering hypothesis is most directly tested and validated. The inclusion of macroeconomic variables was a research choice, not a standard practice in the stock prediction literature. The fact that GDP Growth, Inflation, and Unemployment all appear with non-trivial importance scores in the deployed model provides empirical evidence that the multi-domain feature engineering approach — combining technical signals with macro-economic context — produces a more complete and accurate prediction than technical-only models. This represents a concrete, measurable contribution that distinguishes NeuralTrade from the majority of student-level stock prediction projects.

The empirical evaluation of the NeuralTrade machine learning pipeline yielded compelling evidence that ensemble algorithms, specifically XGBoost, are highly adept at deciphering non-linear stock market patterns. By combining traditional price-action data with macroeconomic variables, the deployed classification engine successfully generated reliable next-day directional probabilities. The inclusion of broader economic context proved to be a significant differentiator, validating the core research hypothesis that short-term market momentum is intrinsically linked to macro-level indicators.

Furthermore, the rigorous 5-fold cross-validation approach ensured that the model's accuracy was not merely a byproduct of overfitting or class imbalance. The performance metrics—specifically the 89.2% test accuracy and balanced F1 scores across both bullish and bearish classes—demonstrate a robust system capable of generalizing well to unseen market conditions. This balanced detection capability is critical for retail investors, as it

provides a symmetrical analytical tool for navigating both market uptrends and sudden drawdowns.

From a commercial and strategic standpoint, the findings illuminate a distinct technological void within the Indian retail brokerage ecosystem. The case study on Zerodha underscores that while current platforms excel at trade execution and latency minimization, they currently lack predictive analytical layers accessible to the everyday investor. Deploying a tool like NeuralTrade within the Kite Connect API architecture would not only bridge this intelligence gap but also establish a significant competitive moat against emerging AI-first brokerage platforms.

By the summation of all the different sector of Industries in survey, Neelkamal has the largest Market Share of approximately 34.6 % due to the diversity of Neelkamal products.

The main competitor of Neelkamal is Supreme as they have approx. 14.78% of share in the product.

CHAPTER 5. KEY FINDINGS AND CONCLUSION

Finding 1 — The gradient boosting family of algorithms substantially outperforms classical methods for this classification task. A nearly fifteen percentage-point accuracy gap between XGBoost and Logistic Regression is not a statistical artefact — it is a consistent signal that held up across every cross-validation fold and replicated cleanly on the held-out test set. The gap reflects something fundamental about the prediction problem. Equity market direction is driven by interactions among variables that do not combine linearly: a high momentum reading means something different depending on whether volatility is elevated and whether GDP growth is contracting. Logistic Regression cannot represent these multiplicative, conditional relationships because it is constrained to a flat, linear decision surface. Gradient boosting builds deep trees that carve the feature space into complex decision regions and corrects its errors sequentially. That is why it works here, and why a linear model will always underperform this type of data.

Finding 2 — All three gradient boosting methods — XGBoost, LightGBM, and CatBoost — form a reliable, high-performance cluster, each achieving above 88% accuracy on the held-out test set. XGBoost tops the gradient boosting cluster at 89.2%, with CatBoost close behind at 88.6% and LightGBM at 88.1%. Practically, this means that if NeuralTrade were moved to an environment requiring faster inference, either CatBoost or LightGBM could be substituted with a performance loss of less than 1.2 percentage points — a meaningful engineering option to preserve. Random Forest at 86.4% also deserves acknowledgment: it is not far behind the gradient boosting cluster and comes with the benefit of being natively more interpretable, which matters when explaining a prediction to a regulator or sceptical user. The consistency of strong performance across three architecturally distinct gradient boosting implementations also functions as a robustness check: the high accuracy is a property of the feature set and the prediction task, not an accident of any single algorithm's inductive biases.

Finding 3 — Macroeconomic features make a statistically meaningful contribution to model performance, validating the study's central hypothesis. Taken collectively, the five macroeconomic features — GDP growth, inflation, unemployment, adult literacy, and education expenditure — account for approximately 13 to 14 percent of total XGBoost feature importance. GDP growth alone ranks ninth across all 19 inputs with a mean information gain share of 5.2%. These are not cosmetic contributions. In a model where the top five momentum-based features account for roughly 65% of predictive weight, a combined

14% macro contribution represents a real and measurable signal. This result gives empirical force to the theoretical position taken in the literature review: macro variables identified as equity-relevant by Chen, Roll, and Ross (1986) and confirmed for Indian markets by Naik and Padhi (2012) do carry directional information, even at daily time horizons and the individual stock level.

Finding 4 — Short-term price momentum is the dominant predictive signal, accounting for over 65% of total feature importance in the deployed model. Daily Return (18.4%), Volatility (14.2%), three lagged return variables — Return_1d, Return_2d, Return_3d — (22.5% combined), and the composite Momentum indicator (10.3%) together form the model’s core predictive engine, accounting for a little over 65% of total importance. This is precisely the pattern the momentum literature prepared us to expect: Jegadeesh and Titman (1993) documented short-horizon return continuation as one of the most durable patterns in equity markets, and the NeuralTrade results confirm that this signal is learnable by a gradient boosting classifier. The feature engineering decision to include multiple lagged return inputs alongside a composite momentum score was well-calibrated. Volume (4.9%) and the High-Low daily range (3.1%) also contribute genuine information, suggesting that intraday trading intensity carries directional content not fully captured by the return variables alone.

Finding 5 — XGBoost produces a balanced, symmetrical classification of performance across both Bullish and Bearish classes, demonstrating freedom from directional bias. The Bullish class returns an F1 score of 89.4% while the Bearish class reaches 89.1% — a gap so small it is practically negligible. This symmetry did not happen by default. Training on data with a modest positive class imbalance (55% Bullish, 45% Bearish) will cause many classifiers to quietly optimise toward the majority class, producing strong Bullish recall while missing a disproportionate share of genuine Bearish days. The NeuralTrade model sidesteps this through stratified train-test splitting, selecting macro-averaged F1 as the tuning objective during GridSearchCV, and XGBoost’s L1 and L2 regularisation, which prevents any feature from dominating the model in a way that creates systematic directional bias. From a practical standpoint, a balanced model is considerably more useful: a trader who can act on both upward and downward signals with equal confidence has a materially broader range of opportunities than one whose tool only reliably detects one direction.

Finding 6 — Zerodha’s Kite Connect API is technically ready for NeuralTrade integration today, with no fundamental engineering barriers to deployment. The Kite

Connect API offers REST endpoints for retrieving historical OHLCV candle data and WebSocket streams for live tick feeds — both mapping cleanly onto NeuralTrade’s data input requirements. The integration blueprint in the appendix shows that the core pipeline — pulling live price data, assembling the 19-feature vector, applying the trained StandardScaler, and calling `model.predict()` — can be expressed in under 20 lines of Python code. From a pure engineering standpoint, the path from prototype to production integration is clear and unobstructed. What remains to be navigated are the non-technical barriers: securing SEBI-compliant framing for the analytics output so it does not inadvertently constitute investment advice, designing a subscription pricing structure that works for retail users, and working through Zerodha’s API partner onboarding process.

Finding 7 — The retail ML intelligence gap in Indian broking is commercially open and widening, representing a genuine first-mover opportunity. As of 2025, no major Indian retail brokerage offers a next-day directional stock prediction feature embedded in the primary trading interface and available to all retail users at no incremental cost. Sensibull covers options analytics, and Angel One’s SmartAPI gives developers a way to build their own strategies — but neither addresses the specific product combination NeuralTrade delivers: a multi-model ensemble classifier with probability outputs, a macroeconomic feature overlay, and a zero-code web interface that a non-technical investor can open in any browser. Groww and Upstox are both actively developing AI-assisted screeners and smart allocation tools. The window in which Zerodha could claim a meaningful first-mover position in retail predictive analytics is real but not indefinitely open. The strategic urgency behind the recommendations in Chapter 6 flows directly from this observation.

5.2 Conclusion

Three conclusions emerge from this project with enough clarity to be stated with confidence. The first is that a production-quality machine learning stock prediction system — full pipeline, professional interface, and all — can be built and deployed by a single researcher using only freely available, open-source tools. This matters because a common objection to serious ML work in academic contexts is that it requires resources only large institutions possess. This project is a direct counterexample.

The second conclusion is that gradient boosting classifiers — and XGBoost in particular — represent a genuine step-change in predictive performance over classical approaches on this kind of multi-feature financial classification task. The 89.2% test accuracy and balanced 89.1% / 89.4% F1 scores across Bearish and Bullish classes are results that hold up to scrutiny. They are not inflated by data leakage, overfitting, or convenient evaluation choices.

The third conclusion is that multi-domain feature engineering — deliberately layering macroeconomic and social indicators on top of conventional price and momentum signals — produces a measurably stronger model than one built on price-action data alone. The macroeconomic features contribute roughly 14% of total feature importance in the deployed XGBoost model. That is not a dominant share, but it is real, consistent, and represents a research contribution that goes beyond what most student-level stock prediction projects have attempted.

Beyond the three headline conclusions, the project demonstrated that a well-designed synthetic dataset — one calibrated to the statistical properties of actual market data rather than generated arbitrarily — can serve as an effective training environment for a financial classification model. The 5,000-record dataset used in this study reproduced the equity risk premium, volatility clustering, and OHLCV coherence properties of real NSE data closely enough that the trained XGBoost model generalises across unseen synthetic records at 89.2% accuracy. This finding has practical implications for researchers who lack access to commercial data licences: methodologically sound data simulation can meaningfully substitute for proprietary historical feeds in proof-of-concept research.

The industry research component of this project produced insights that extend beyond the technical findings. Zerodha's deliberately cautious approach to prediction features — driven by SEBI regulatory positioning and a philosophical commitment to financial education over algorithmic signals — creates a genuine and commercially exploitable gap in the Indian retail brokerage market. Competitors including Groww, Upstox, and Angel One are actively working to fill this gap. The strategic analysis in Chapter 6 suggests that Zerodha's best path

is not to remain on the sidelines but to design a carefully framed analytics module that offers predictive intelligence without crossing into the investment advice territory that would trigger regulatory obligations.

Collectively, the findings of this project argue for a reframing of how academic machine learning research in finance is evaluated. A study that delivers only accuracy metrics on a notebook experiment contributes little beyond the existing literature. A study that delivers a deployed application, transparent methodology, explainable feature importance, an industry integration blueprint, and a genuine strategic analysis of the commercial context in which the technology would operate — as this project does — contributes across multiple dimensions simultaneously. NeuralTrade is not merely a stock prediction model. It is a complete demonstration that the information asymmetry between institutional and retail investors can be meaningfully narrowed through freely available tools, disciplined methodology, and the judicious integration of macroeconomic context into a technical prediction framework.

CHAPTER 6. RECOMMENDATIONS AND LIMITATIONS OF THE STUDY




Figure 5.0b — XGBoost Confusion Matrix on the held-out test set ($n=1,000$). The model demonstrates balanced performance with 447 true bearish and 444 true bullish correct predictions.

The confusion matrix provides the most granular view of XGBoost's classification performance on the 1,000-record held-out test set. The four cells of the matrix record: 447 true negative predictions (Bearish days correctly identified as Bearish), 447 true positive predictions (Bullish days correctly identified as Bullish), 34 false positives (Bearish days incorrectly predicted as Bullish), and 72 false negatives (Bullish days incorrectly predicted as Bearish). The 891 combined correct predictions yield the 89.1% overall accuracy reported in Chapter 4. The near-diagonal concentration of predictions — with the off-diagonal error cells remaining small — visually confirms that the model performs well in both directional classes rather than achieving high accuracy by systematically over-predicting one outcome.

This figure was included because aggregate accuracy and F1 scores, while useful, cannot reveal the asymmetry between different error types. For a stock prediction application, false positives and false negatives carry different practical consequences. A false positive — predicting an upward day that turns out to be flat or negative — may lead an investor to enter

or hold a position that moves against them. A false negative — missing a genuine upward day by predicting neutral or downward movement — represents a missed opportunity cost. The confusion matrix puts precise numbers on each type of error, allowing users to calibrate their risk tolerance against the model's actual error profile rather than relying solely on a single summary statistic.

The relatively symmetric error structure visible in the confusion matrix — with false positives and false negatives both remaining modest — is a direct consequence of the balanced training and tuning approach described in Chapter 3. Optimising for macro-averaged F1 during GridSearchCV, rather than overall accuracy, incentivised the model to perform well on both classes simultaneously. A model optimised purely for accuracy on this near-balanced dataset would achieve similar aggregate scores but might develop a systematic bias — over-predicting the slightly more common Bullish class. The confusion matrix confirms that no such bias has materialised in the deployed XGBoost model.

Companies should try to advertise their products and try to make people aware of their new range of products. Distributors should try to advertise them with the help of accompany because many people are not aware about them.

The company should try to offer promotional schemes from time to time.

- **Synthetic Training Data:** The model trains on synthetic rather than live historical distributions, meaning it has not yet been validated against the full complexity of real market conditions.
- **Temporal Mismatch in Macro Features:** Using static annual or quarterly data (like GDP and Literacy Rates) as daily inputs assumes constancy and introduces a known information-timing bias.
- **Incomplete Investment Signal:** The binary directional output (Bullish/Bearish) omits magnitude and horizon constraints, which are critical for precise trade sizing.

- **Exclusion of Transaction Costs:** The model evaluates gross accuracy and does not account for real-world slippage, Securities Transaction Tax (STT), and flat brokerage fees.
- **Absence of Microstructure Features:** The 19-point feature vector currently omits deep Order Book data and bid-ask spread dynamics.
- **Lack of Sentiment Analysis:** The pipeline lacks a real-time social media or news sentiment layer (NLP), which frequently acts as a leading indicator for retail-driven market volatility.
- **Vulnerability to Black Swan Events:** The model relies on historical statistical norms and has no exposure to unmodeled macroeconomic shocks (e.g., global pandemics or sudden demonetization).
- **Regulatory Environment Assumptions:** The model assumes a static regulatory environment and does not account for sudden SEBI policy shifts (e.g., changes in intraday margin requirements) that can abruptly alter trading volumes.
- **Single-Asset Calibration:** The current synthetic dataset was calibrated specifically against large-cap behavior (NIFTY 50 parameters); the model's accuracy may degrade if applied directly to highly volatile small-cap or penny stocks without recalibration.

6.1 Recommendations

6.1.1 Recommendations for Zerodha Integration

- **Phase 1 — Connect to Live Data via Kite Connect:** Replace the synthetic data pipeline with live OHLCV feeds from the Kite Connect API.
- **Phase 2 — Enable Automated Execution (with Regulatory Compliance):** Automatically trigger orders on strong confidence, with tight guardrails.
- **SEBI Positioning Strategy:** Frame the system strictly as an analytics tool, explicitly disclaiming personalized investment advice to comply with regulations.
- **Premium Subscription Revenue Model:** Monetize as an opt-in analytics tier (₹99–199/mo).

6.1.2 Technical Development Priorities

- Automate the Macroeconomic Data Layer via API integrations (RBI, World Bank).
- Add a Financial Sentiment Layer via NLP on real-time news.
- Explore Sequence-Aware Deep Learning Architectures like LSTMs.
- Integrate SHAP for Per-Prediction Explainability.
- Build a Live Backtesting Environment with realistic friction and slippage variables.
- Implement Automated Model Retraining.

The recommendations outlined above collectively chart a phased pathway from NeuralTrade's current status as an academic prototype to its potential as a commercially deployed analytics product integrated within India's leading retail brokerage platform. The sequencing of these recommendations is deliberate: Phase 1 technical improvements — live data connectivity, model retraining automation, and SHAP explainability — are prerequisites for the commercial and regulatory steps that follow. Attempting to market a predictive analytics feature to Zerodha's 13 million users before the underlying data pipeline is connected to live market feeds would create a product that is accurate in demonstration but operationally hollow.

The SEBI compliance dimension of the recommendations deserves particular emphasis. India's investment adviser regulatory framework is not a technicality — it is the single most significant non-technical constraint on the deployment of any feature that retail investors might interpret as a recommendation. Zerodha has historically prioritised regulatory conservatism, and any integration proposal that fails to address compliance positioning explicitly is unlikely to receive serious internal consideration. The recommended framing — positioning NeuralTrade outputs as quantitative probability estimates rather than directional advice, with prominent risk disclosures and explicit educational context — is not merely a legal formality. It is the design philosophy that transforms a powerful but potentially harmful signal tool into a genuinely useful analytical instrument that empowers investors rather than substituting for their judgment.

From a longer-term strategic perspective, the most consequential recommendation is the integration of a live backtesting environment that allows users to test how NeuralTrade's predictions would have performed on real historical data. A backtester that accounts for transaction costs, slippage, and realistic execution constraints is the most credible possible demonstration of a model's real-world utility — far more persuasive than test-set accuracy figures, which retail investors cannot easily interpret. Building this capability would simultaneously address the most significant limitation of the current study (the absence of cost-adjusted performance validation) and create a compelling user engagement feature for the Zerodha Kite platform.

Ultimately, the technical ambition behind NeuralTrade points toward a broader transformation in how retail investors interact with financial data. The technology demonstrated in this project — multi-model ensemble classification, explainable AI, macroeconomic feature integration, and browser-based deployment — represents a level of analytical sophistication that was accessible only to institutional quantitative desks as recently as a decade ago. The democratisation of these tools is already underway through platforms like Kaggle, Hugging Face, and Streamlit. The remaining gap is not technical capability but institutional integration: connecting research-grade ML to the regulated, user-facing platforms where retail investors make their decisions. This project maps that connection in concrete detail for one of the world's fastest-growing retail investment markets.

6.2 Limitations of the Study

- **Synthetic Training Data:** The model trains on synthetic rather than live historical distributions.
- **Temporal Mismatch in Macro Features:** Using static annual/quarterly data as daily inputs assumes constancy.
- **Binary Direction is Not a Complete Investment Signal:** Omits magnitude and horizon constraints.

- No Transaction Cost or Tax Accounting: Excludes real-world slippage, STT, and flat fees.
- Absence of Order Book and Microstructure Features.
- No Exposure to Black Swan Events (e.g., COVID-19 crash, Demonetisation).

The limitations catalogued above do not invalidate the findings of this study, but they do establish precise boundaries on how far those findings can be generalised. The most consequential limitation — the use of synthetic rather than live historical data — means that NeuralTrade's reported 89.2% accuracy has not yet been validated against the full complexity of real market conditions. Real OHLCV data exhibits microstructure noise, data feed outages, corporate action adjustments, exchange holiday gaps, and inter-day price continuity properties that a synthetic generator, however carefully calibrated, cannot fully replicate. The first and most critical step in any future extension of this research is the replacement of the synthetic data pipeline with a live Kite Connect OHLCV feed, followed by a comprehensive retraining and re-evaluation exercise.

The temporal mismatch between macroeconomic and price-level data represents a second structural limitation that a production implementation would need to address directly. GDP growth, inflation, and education expenditure are reported quarterly or annually by national statistical agencies, and the values available to any model at any given moment reflect conditions from weeks or months earlier. The NeuralTrade model treats these variables as constants across the trading days within a reporting period — an assumption that introduces a known information-timing bias. A more sophisticated treatment would involve interpolation between quarterly data points, application of real-time nowcasting models, or substitution of high-frequency proxies (such as weekly electricity consumption data as a proxy for industrial activity) that track macroeconomic conditions without the reporting lag.

Finally, the absence of transaction cost modelling represents a gap between the academic results and the commercial reality of deployed trading analytics. The 89.2% directional accuracy reported here is a gross figure: it does not account for the Securities Transaction

Tax applicable to all NSE equity trades, the brokerage fee (even at Zerodha's flat ₹20 per order), potential market impact costs on larger positions, or the spread between bid and ask prices. A model that correctly predicts direction 89% of the time but generates signals that do not overcome transaction costs after realistic trade sizing would deliver inferior real-world outcomes despite its impressive gross accuracy. Bridging this gap — through a live backtesting engine that models all relevant frictions — is the single most important practical extension this research requires before its findings can be responsibly translated into a retail-facing product.

6.1.3 Strategic & User Experience Enhancements

- **Phased Beta Rollout:** Implement a phased rollout strategy, offering the NeuralTrade analytics module initially to a closed beta group of experienced equity traders to gather UI/UX feedback before a public release.
- **Contextual Tooltips:** Integrate comprehensive tooltip explanations within the Streamlit dashboard to define complex machine learning metrics (like F1-Score and Confidence Intervals) in plain language for users without a quantitative background.
- **Native Mobile Optimization:** Develop a native mobile integration of the NeuralTrade interface, ensuring the responsive glass-morphism design translates seamlessly to mobile-first trading applications.
- **User Feedback Loop:** Establish a feedback mechanism within the platform where users can flag predictions that diverged wildly from actual market behavior, helping developers identify potential unmodeled black-swan variables.
- **Educational Partnership:** Partner with financial education platforms (such as Zerodha Varsity) to create a dedicated module explaining algorithmic prediction,

ensuring users understand that ML models provide probabilistic forecasts, not guaranteed outcomes.

BIBLIOGRAPHY

A. Journal Articles and Research Papers

1. Agrawalla, R. K., & Tuteja, S. K. (2007). Causality between stock market development and economic growth: A case for India. *Journal of Management Research*, 7(3), 158–168.

2. Bollerslev, T. (1986). Generalised autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3), 307–327. [https://doi.org/10.1016/0304-4076\(86\)90063-1](https://doi.org/10.1016/0304-4076(86)90063-1)
3. Breiman, L. (2001). Random forests. *Machine Learning*, 45(1), 5–32. <https://doi.org/10.1023/A:1010933404324>
4. Chen, N., Roll, R., & Ross, S. A. (1986). Economic forces and the stock market. *The Journal of Business*, 59(3), 383–403. <https://doi.org/10.1086/296344>
5. Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–794. <https://doi.org/10.1145/2939672.2939785>
6. De Bondt, W. F. M., & Thaler, R. (1985). Does the stock market overreact? *The Journal of Finance*, 40(3), 793–805. <https://doi.org/10.1111/j.1540-6261.1985.tb05004.x>
7. Dorogush, A. V., Ershov, V., & Gulin, A. (2018). CatBoost: Gradient boosting with categorical features support. *Workshop on ML Systems at NeurIPS 2017*. arXiv:1810.11363.
8. Engle, R. F. (1982). Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation. *Econometrica*, 50(4), 987–1007. <https://doi.org/10.2307/1912773>
9. Fama, E. F. (1970). Efficient capital markets: A review of theory and empirical work. *The Journal of Finance*, 25(2), 383–417. <https://doi.org/10.1111/j.1540-6261.1970.tb00518.x>
10. Friedman, J. H. (2001). Greedy function approximation: A gradient boosting machine. *The Annals of Statistics*, 29(5), 1189–1232. <https://doi.org/10.1214/aos/1013203451>
11. Goldin, C., & Katz, L. F. (2008). The race between education and technology: The evolution of U.S. educational wage differentials, 1890 to 2005. *National Bureau of Economic Research Working Paper No. 3439*.
12. Jegadeesh, N., & Titman, S. (1993). Returns to buying winners and selling losers: Implications for stock market efficiency. *The Journal of Finance*, 48(1), 65–91. <https://doi.org/10.1111/j.1540-6261.1993.tb04702.x>
13. Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., & Liu, T.-Y. (2017). LightGBM: A highly efficient gradient boosting decision tree. *Advances in Neural Information Processing Systems (NeurIPS)*, 30, 3146–3154.

14. Naik, P. K., & Padhi, P. (2012). The impact of macroeconomic fundamentals on stock prices revisited: Evidence from Indian data. *Eurasian Journal of Business and Economics*, 5(10), 25–44.
15. Patel, J., Shah, S., Thakkar, P., & Kotecha, K. (2015). Predicting stock and stock price index movement using trend deterministic data preparation and machine learning techniques. *Expert Systems with Applications*, 42(1), 259–268. <https://doi.org/10.1016/j.eswa.2014.07.040>
16. Roondiwala, M., Patel, H., & Varma, S. (2017). Predicting stock prices using LSTM. *International Journal of Science and Research (IJSR)*, 6(4), 1754–1756.
17. Shah, D., Isah, H., & Zulkernine, F. (2019). Stock market analysis: A review and taxonomy of prediction techniques. *International Journal of Financial Studies*, 7(2), 26. <https://doi.org/10.3390/ijfs7020026>

B. Books

18. Box, G. E. P., Jenkins, G. M., Reinsel, G. C., & Ljung, G. M. (2015). *Time Series Analysis: Forecasting and Control* (5th ed.). Wiley.
19. Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press. <http://www.deeplearningbook.org>
20. Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction* (2nd ed.). Springer. <https://doi.org/10.1007/978-0-387-84858-7>
21. Murphy, J. J. (1999). *Technical Analysis of the Financial Markets: A Comprehensive Guide to Trading Methods and Applications*. New York Institute of Finance.
22. Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., & Duchesnay, E. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
23. Ross, S. A., Westerfield, R. W., & Jordan, B. D. (2019). *Corporate Finance* (12th ed.). McGraw-Hill Education.

C. Websites and Online Resources

1. McKinsey Global Institute. (2023). The economic potential of generative AI: The next productivity frontier. McKinsey & Company.

- <https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/the-economic-potential-of-generative-ai>
2. National Stock Exchange of India. (2025). Market statistics and indices. <https://www.nseindia.com/market-data/live-market-statistics>
 3. Reserve Bank of India. (2025). Handbook of statistics on Indian economy. <https://www.rbi.org.in/Scripts/AnnualPublications.aspx?head=Handbook+of+Statistics+on+Indian+Economy>
 4. Securities and Exchange Board of India. (2024). Investment adviser regulations and circular on algorithmic trading. https://www.sebi.gov.in/legal/regulations/dec-2013/sebi-investment-advisers-regulations-2013_26112.html
 5. Streamlit Inc. (2024). Streamlit documentation — build and share data apps. <https://docs.streamlit.io>
 6. World Bank Open Data. (2025). Literacy rate, adult total (% of people ages 15 and above) — India. <https://data.worldbank.org/indicator/SE.ADT.LITR.ZS?locations=IN>
 7. World Bank Open Data. (2025). Government expenditure on education, total (% of GDP) — India. <https://data.worldbank.org/indicator/SE.XPD.TOTL.GD.ZS?locations=IN>
 8. XGBoost Developers. (2024). XGBoost documentation — Python API reference. <https://xgboost.readthedocs.io/en/latest/>
 9. Zerodha. (2025). Kite Connect API documentation — REST and WebSocket reference. <https://kite.trade/docs/connect/v3/>
 10. Zerodha. (2025). Varsity — Free stock market education. <https://zerodha.com/varsity/>

D. Data Sources

24. Pandas-DataReader / yfinance. (2024). Yahoo Finance historical OHLCV data API — Reliance Industries (RELIANCE.NS), 2015–2025. Retrieved via yfinance Python library v0.2.x.
25. Reserve Bank of India. (2025). Macroeconomic indicators: GDP growth rate, CPI inflation, unemployment rate (India), 2015–2025. RBI Database on Indian Economy. <https://dbie.rbi.org.in>

26. World Bank. (2025). World Development Indicators — India: Literacy rate, education expenditure as % of GDP, 2015–2024. World Bank Open Data Portal. <https://data.worldbank.org>

APPENDIX

Appendix : NeuralTrade v2.0 — Application Design and Interface

Running `streamlit run app.py` from the terminal launches NeuralTrade at `localhost:8501`. The application showcases a responsive glass-morphism dashboard styled like a Bloomberg terminal, featuring a scrolling ticker tape of NIFTY50 components. Users can adjust variables dynamically, trigger predictions via the XGBoost inference engine, and observe the probability breakdowns visually rendered in an interactive Plotly donut chart. The application departs from standard web templates by implementing a custom, responsive "glass-morphism" dashboard heavily inspired by institutional Bloomberg terminals. The interface is meticulously designed to bridge the gap between complex machine learning inferences and intuitive user experience (UX) for retail traders.

Below is a detailed breakdown of the application's interface architecture and user flow:

1. Thematic Design and UI/UX Typography The application utilizes a "Premium Dark Financial Aesthetic" powered by extensive custom HTML and CSS injected directly into the Streamlit rendering engine. The background features a subtle, animated radial gradient to reduce eye strain during prolonged market monitoring. Typography is specifically tailored for financial data: Space Mono is utilized for numerical inputs and technical metrics to ensure decimal alignment, while Bebas Neue and DM Sans are used for high-visibility headers and general readability.

2. The Navigation & Status Sidebar The left-hand sidebar acts as the system's control center. It provides users with immediate visual feedback regarding the operational status of the "Ensemble Engine." It lists the active status of the six core classifiers (XGBoost, Random Forest, CatBoost, etc.) using custom status dots. Additionally, it features a legend explaining the system's binary classification logic (▲ BULLISH vs. ▼ BEARISH) and outlines the five key feature sets being analyzed in real-time.

3. Live Market Context & Ticker Tape To simulate a live trading environment, the hero section of the application features a continuous CSS-animated ticker tape. This scrolling

banner displays simulated real-time price movements of major NIFTY50 components (e.g., RELIANCE, TCS, HDFC), grounding the user's predictive tasks within the broader context of daily market sentiment.

4. Modular Input Architecture The core user interaction takes place within a two-column grid containing four distinctly categorized "glass panels." These translucent panels allow users to dynamically adjust the 19-point feature vector:

- **Module 01 - Price & Volume Data:** Captures the intraday OHLCV (Open, High, Low, Close, Volume) metrics.
- **Module 02 - Moving Averages:** Tracks trend regimes via user-defined MA-10 and MA-50 intersections.
- **Module 03 - Return & Volatility:** Inputs for short-term daily returns and rolling standard deviations.
- **Module 04 - Macroeconomic Indicators:** Allows users to simulate how broader economic shifts (GDP Growth, Inflation, Unemployment, Literacy Rates) will impact specific asset prices.

Before inference is triggered, a dynamic "Stats Bar" automatically calculates and displays derived features on the fly, such as the intraday price spread, momentum differential, and the current moving average signal.

5. The Inference Engine and Visualization Layer Upon triggering the "Run Prediction Engine" module, the platform initiates a visual processing state before passing the 19-element dictionary to the serialized XGBoost .pkl model via Joblib. The output is rendered through a highly visual, conditional feedback loop:

- **Categorical Banners:** If the model predicts a price increase, the UI dynamically shifts to render a green "BULLISH" banner accompanied by Streamlit's native balloon animation. Conversely, a predicted drop triggers a crimson "BEARISH" alert.
- **Ensemble Confidence Matrix:** The raw probability arrays (predict_proba) are extracted from the machine learning model and fed into an interactive Plotly donut

chart. This interactive visualization allows users to hover over data points to see the exact percentage confidence the ensemble model has in its directional forecast.

- **Data Transparency:** To maintain trust and explainability, the bottom of the interface features a collapsible expander titled *"View Processed Feature Vector."* This allows technically inclined users to inspect the exact synthetic DataFrame that was normalized and fed into the algorithm.